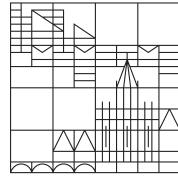


Universität
Konstanz



Beer-Paths

Bachelor Project

Course of Studies Informatik

University Konstanz

Jakob Sanowski

Submitted on: 01.01.1980

Student ID, Course: 7654321, TINF22B2

Supervisor at Uni KN: Prof. Dr. Sabine Storandt

Contents

1	Introduction	1
2	Preliminaries	1
2.1	Problem Definition	2
3	Algorithms & Implementation	3
3.1	Double Dijkstra	3
3.2	Parallel Dual Dijkstra	3
3.3	Beer-Path Oracle	4
3.3.1	In-Path Property	5
3.3.2	R*-Tree	7
4	Experimental Evaluation	8
4.1	Dataset	9
4.1.1	Comparitive Experiments	9
4.1.2	Baseline Approach	9
4.2	In-Path Oracle	9
4.2.1	Varying Detour Limits	9
4.3	Throughput Experiment	10
5	Conclusions and Future Work	11
	Bibliography	12

1 Introduction

In graph theory and computer science, the beer-path problem presents a unique challenge that extends traditional shortest path queries by introducing the necessity to traverse specific vertices, known as “beer vertices.” This problem is particularly relevant in scenarios where paths must include certain checkpoints or resources, analogous to visiting a “beer store” in a network of roads. The beer-path oracle is a specialized data structure designed to efficiently answer queries related to beer paths, providing all beer vertices which in-path for any two vertices.

This report delves into the performance of a beer-path oracle, exploring its efficiency, scalability, and practical applications. We begin by outlining the theoretical foundations of the beer-path problem, highlighting its significance in various computational contexts such as network routing and logistics. The core of this report focuses on the implementation details of the beer-path oracle, including the algorithms and data structures employed to achieve optimal query times.

We present a comprehensive performance analysis, evaluating the oracle’s response time and memory usage across different types of graphs. Through empirical testing, we demonstrate the oracle’s ability to handle large-scale (lol) graphs and discuss the trade-offs between preprocessing time and query efficiency. Furthermore, we compare our beer-path oracle with a double dijkstra approach, underscoring its advantages and potential areas for improvement.

The findings of this report contribute to the ongoing research in graph algorithms and data structures, offering insights into the development of efficient pathfinding techniques under constrained conditions. By understanding the performance characteristics of the beer-path oracle, we aim to provide a robust framework for future advancements in this field.

2 Preliminaries

In this section we will establish some preliminary concepts and describe the problem itself. Most of the definitions are taken from D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1].

Definition 1**Shortest Distance**

Given source s and destination t nodes, $d_N(s, t)$ denotes the shortest distance between s and t . $d_N(s, t)$ is obtained by summing over the edge weights along the shortest path between s and t .

Definition 2**Detour**

Given source s and destination t nodes, let $\pi(s, t)$ denote a simple path that is not necessarily the shortest. The detour of such a path is the difference in the network distance along $\pi(s, t)$ compared to $d_N(s, t)$. Furthermore, it is fairly trivial to see that the detour of any path is greater or equal to zero.

Definition 3**Detour Bound**

The detours are bounded by a fraction ε such that their total distance does not exceed $\varepsilon * d_N(s, t)$. For example, if $\varepsilon = 0.1$ a detour can be up 10% longer than the shortest path.

Definition 4**In-Path POI**

A POI is said to be *in-path* if there exists a detour bounded by ε which passes through said POI.

2.1 Problem Definition

We are given a road network G , set P of m POIs, and a detour bound ε . A driver travels from source s and destination t , we want to find the set of pois in p that are “in-path” under the conditions specified.

Now that we have provided our problem statement we can also discuss the limitations of our approach. First, we are interested in retrieving all the POIs that satisfy the detour constraints. Our focus here is maximising the throughput where one can answer millions of in-path queries a second using a single machine. Our solution is not geared towards a driver that wants to visit multiple POIs yet stay within the detour bound. In our model, the expectation is that the user is presented with the POI choices and may choose one of the in-path POI to visit. Such examples include coffee shops, restaurants, gas stations, vaccination clinics, etc. The driver is unlikely to visit another POI of the same kind. Our work is in the context of the placement of relevant POIs on a map as an

opportunistic service where speed is of the essence, so, the composition of complex trips that include visiting multiple POIs is not the focus of this work.

3 Algorithms & Implementation

In this section we will look at the algorithms we want to compare in this report. The first algorithm is a double dijkstra exploring from the start and target towards the POIs. The second algorihm is a parallel version of the dual Dijkstra [1]. The last algorithm uses a in-path oracle [1] for faster query times.

3.1 Double Dijkstra

The double Dijkstra is a Dijkstra variant for finding detours passing through one $p \in P$. We use two seperate instances of Dijkstra starting from the start s and end t node respectively. The input for both instances are all POIs from P and t for the instance starting from s . We combine the result of both instances by adding the costs from both instances for every $p \in P$ together. It is important to note for the instance starting from t we traverse the edges backwards.

3.2 Parallel Dual Dijkstra

D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1] proposed the dual Dijkstra algorithm for finding POIs within a specified detour tolerance limit ε which we developed a parallel version of. In order to parallelise the algorithm we run two Dijkstra at the same time starting from the source s and destination t similar to the double Dijkstra.

Algorithm 1 describes the algorithm of both instances. Each instance uses its own a priority queue Q over the distance to it's respective start node. Every node n additionally holds the distance to the start and a label which can be accessed with the functions $d(n)$ and $l(n)$.

At the core of this algorithm is the shared data structure **VISITED**. This data structure holds all nodes visited by both Dijkstra instances together with a label indicating which instance found the node and the distance to the start node s or t respectively. The key of this algorithm is in Line 11 where we add the two distances together. If this node $n \in P$ we mark it as **POI** so it gets added to the result.

Algorithm 1: Dual Dijkstra

Data:**Result:**

```

1   while  $!Q.empty()$   $n := Q.front()$   $d(n) \leq d_N$  do
2     if  $l(n) == \mathbb{P}\mathbb{O}\mathbb{I}$  then
3       |  $result.add()$ 
4       | continue
5     end
6     if  $\text{VISITED}(n, l(n))$  do
7       | continue
8     end
9     if  $n_r := \text{VISITED}(n, l(n).inverse())$  do
10      |  $d' := d(n) + d(n_r)$ 
11      |  $n.distance(d')$ 
12      |  $d_N := \min(d_N, d' * (1 + \varepsilon))$ 
13      | if  $n \in P$  do
14        | |  $Q.insert(n.label(\mathbb{P}\mathbb{O}\mathbb{I}))$ 
15      | end
16    end
17     $\text{VISITED}.insert(n)$ 
18    for neighbour  $v_i$  of  $n$  do
19      |  $Q.insert(v_i.label(l(n)))$ 
20    end
21  end
22  return  $result$ 

```

3.3 Beer-Path Oracle

The beer-path oracle proposed by D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1] aims to reduce query times using precomputed results. It uses the *spatial coherence* [2] property in road networks which observes similar characteristics for nodes spatially adjacent to each other. Or more precisely the coherence between shortest paths and distances between nodes and their spatial locations [2], [3]. We know for

a set of source nodes A and destination nodes B they might share the same shortest paths if A and B are sufficiently far apart and the nodes contained in A and B are close together. This enables determining if a POI is in-path with respect to this group of nodes opposed to single pairs of nodes.

3.3.1 In-Path Property

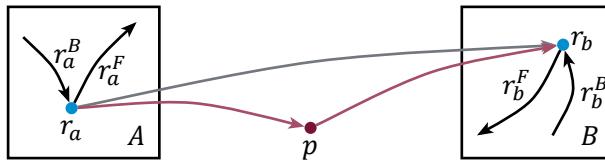


Figure 1 — Whether p is in-path with respect to all sources in A to destinations in B .

In order to define the *in-path* property for a set of source nodes A and a set of destination nodes B these sets are restricted to be inside of a bounding box containing all nodes. Let a_r be a randomly chosen representative source node in A and b_r a representative destination node in B . Let p be the POI we want to determine as in-path with respect to the block-pair (A, B) if all shortest-paths from all sources in A to all destinations in B are in-path to p .

We start by defining r_a^F as the forward radius of a given block A denoting the farthest distance from a_r to any node. Similarly, r_a^B defines the backwards radius denotes the farthest distance of any node to a_r . We also define the forward and backwards radius for any block B as r_b^F and r_b^B respectively (see Figure 1). The following lemmas define bounds for the shortest and longest shortest-paths for all shortest-paths from A to B .

Lemma 1

Shortest Shortest Path

Any shortest path between A and B has a length equal to or greater than

$$d_N(a_r, b_r) - (r_a^F + r_b^B).$$

PROOF: Let s and t be an arbitrary source and destination with $d_N(s, t) < d_N(a_r, b_r)$.

Now one can consider the path $a_r \rightarrow s \rightarrow t \rightarrow b_r$. Note that $a_r \rightarrow s$ is bounded by r_a^B and $t \rightarrow b_r$ is bounded by r_b^F . Following this $d_N(s, t) \geq d_N(a_r, b_r) - (r_a^B + r_b^F)$ has to hold. If $d_N(s, t) < d_N(a_r, b_r) - (r_a^B + r_b^F)$ then $d_N(a_r, b_r)$ would not be the shortest distance between a_r and b_r because $d_N(a_r, s) \leq r_a^B$ and $d_N(t, b_r) \leq r_b^F$ which leads to $d_N(a_r, b_r) < d_N(a_r, b_r) - (r_a^B + r_b^F) + (r_a^B + r_b^F) = d_N(a_r, b_r)$ which is a contradiction. ■

Lemma 2**Longest Shortest Path**

Any shortest path between A and B has a length of at most

$$d_N(a_r, b_r) + (r_a^B + r_b^F)$$

PROOF: Let s and t be an arbitrary source and destination. Then one can define the following path: $s \rightarrow a_r \rightarrow b_r \rightarrow t$. This path is bound by $d_N(a_r, b_r) + (r_a^B + r_b^F)$.

■

Lemma 3**In-Path Property**

A block-pair (A, B) is in-path if the following condition is satisfied and $d_N(a_r, b_r) - (r_a^F + r_b^B) > 0$:

$$\frac{r_a^B + d_N(a_r, p) + d_N(p, b_r) + r_b^F}{d_N(a_r, b_r) - (r_a^F + r_b^B)} - 1 \leq \varepsilon$$

PROOF: For any given node s, t in A, B , respectively, $d_N(s, t)$ is at least $d_N(a_r, b_r) - (r_a^F + r_b^B)$ (see Lemma 1). Considering the path $s \rightarrow a_r \rightarrow p \rightarrow b_r \rightarrow t$ it has a length of at most $r_a^B + d_N(a_r, p) + d_N(p, b_r) + r_b^F$. If p is in-path to $a_r \rightarrow b_r$ then we get the following inequality in order for all possible paths in A, B to be in-path:

$$\begin{aligned} r_a^B + d_N(a_r, p) + d_N(p, b_r) + r_b^F &\leq (d_N(a_r, b_r) - (r_a^F + r_b^B)) \cdot (1 + \varepsilon) \\ \frac{r_a^B + d_N(a_r, p) + d_N(p, b_r) + r_b^F}{d_N(a_r, b_r) - (r_a^F + r_b^B)} - 1 &\leq \varepsilon \end{aligned}$$

■

Note that the condition $d_N(a_r, b_r) - (r_a^F + r_b^B) > 0$ is omitted by D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1] but is necessary because $d_N(a_r, b_r)$ can be 0 in which case $d_N(a_r, b_r) - (r_a^F, r_b^B) < 0$ and thus the condition would suddenly be satisfied if $d_N(a_r, b_r)$ is smaller than some specific value.

Lemma 4**Not In-Path Property**

A block pair (A, B) is not in-path if the following condition is satisfied:

$$\frac{d_N(a_r, p) + d_N(p, b_r) - (r_a^B + r_b^F)}{d_N(a_r, b_r) + (r_a^B + r_b^F)} - 1 \geq \varepsilon$$

PROOF: For any given node s, t in A, B , respectively, $d_N(s, t)$ is at most $d_N(a_r, b_r) + (r_a^B + r_b^F)$ (see Lemma 2). Considering the path $s \rightarrow a_r \rightarrow p \rightarrow b_r \rightarrow t$ it has a length of at least $d_N(a_r, p) + d_N(p, b_r) - (r_a^B + r_b^F)$. We get the following inequality in order for all possible paths in A, B to not be *in-path* to p :

$$\begin{aligned} d_N(a_r, p) + d_N(p, b_r) - (r_a^B + r_b^F) &\geq (d_N(a_r, b_r) + (r_a^B + r_b^F)) \cdot (1 + \varepsilon) \\ \frac{d_N(a_r, p) + d_N(p, b_r) - (r_a^B + r_b^F)}{d_N(a_r, b_r) + (r_a^B + r_b^F)} - 1 &\geq \varepsilon \end{aligned}$$

■

Algorithm 2 – In-Path Oracle for a given POI

```

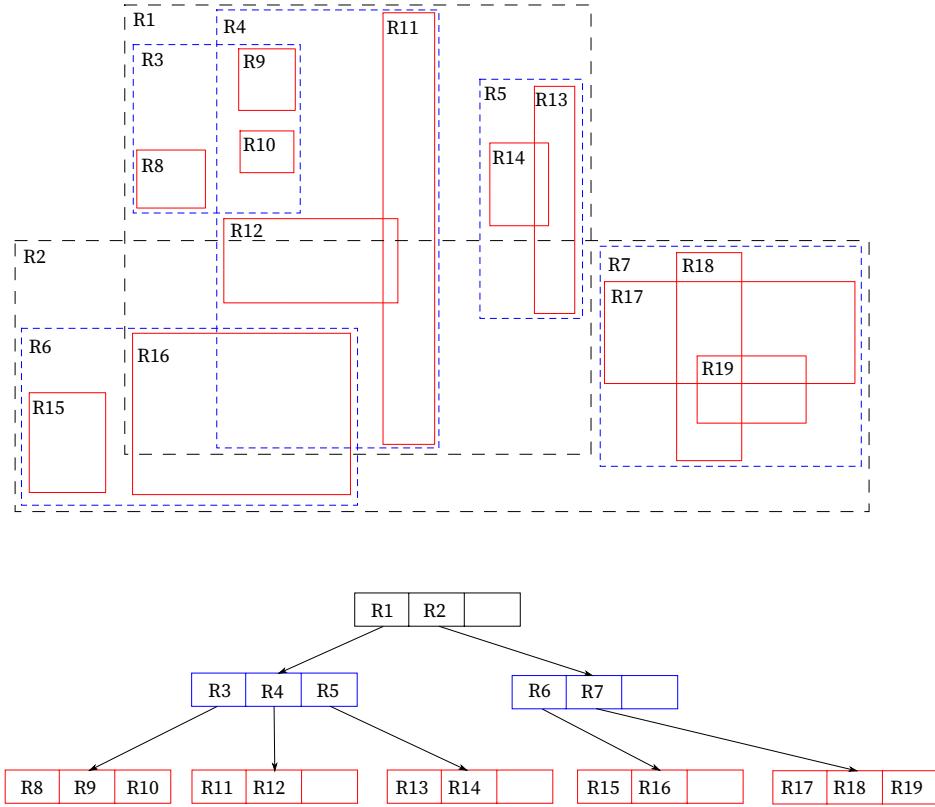
1:  $R \leftarrow$  root block of the road network
2:  $result \leftarrow \emptyset$ 
3:  $Q \leftarrow \{R, R\}$ 
4: while  $!Q.empty()$  do
5:    $(A, B) \leftarrow Q.pop\_front()$ 
6:    $s, t \leftarrow$  random node from  $A, B$ , respectively
7:    $values \leftarrow$  Compute  $d_n(s, t), d_N(s, p), d_N(p, t), r_a^F, r_a^B, r_b^F, r_b^B$ 
8:   if  $values.in-path()$  then
9:      $result.add((A, B))$ 
10:  end
11:  if  $values.not-in-path()$  then
12:    continue
13:  end
14:  Subdivide  $A$  and  $B$  into 4 children blocks. Discard empty children blocks.
15:  Insert all children blocks into  $Q$ 
16: end

```

3.3.2 R*-Tree

In order to get fast query times we used an *R*-Tree* [4] for storing the oracle. The *R*-Tree* is a variant of the *R-Tree* [5] which tries to minimize overlap.

The idea behind *R-Trees* is to group nearby objects into rectangles and in turn store them in a tree similar to a *B-Tree* (see Figure 2). Also like in a *B-Tree* the data is organized into pages of a fixed size. This enables search similarly to a *B-Tree* recursively searching through all nodes which bounding boxes are overlapping with the search area.

Figure 2 — *R*-Tree for 2D rectangles with a page size of 3

The performance of an *R*-Tree greatly depends on the overlap of the bounding boxes in the tree. Generally less overlap leads to better performance. For this reason the insertion strategy is crucial for achieving good performance. *R*^{*}-Trees try to minimize the overlap by employing insertion strategies which take this into account. This improves pruning performance, allowing exclusion of whole pages from search more often. The key for achieving this is based on the observation that *R*-Trees are highly susceptible to the order in which their entries are inserted. For this reason the *R*^{*}-Tree performs reinsertion of entries to “find” a better suited place in the tree.

In the case of a node overflowing a portion of its entries are removed and reinserted into tree. To avoid infinite reinsertion, this may only be performed once per level of the tree.

4 Experimental Evaluation

The experiments were performed on an AMD Ryzen 5 5600X with 6 cores and 12 threads at 4.651 GHz and 16 GB of RAM.

4.1 Dataset

The road networks used for evaluation were obtained from OpenStreetMap and sanitized of foot-paths to only include one edge per street. We used two datasets in our evaluation, Konstanz with 2282 nodes and 4377 edges and San Francisco with 95092 nodes and 172256 edges. The weight of each directed edge denotes the travel distance between two nodes. Note that *chains* (or *ways*) are not simplified.

4.1.1 Comparative Experiments

We used the dual Dijkstra as a baseline for comparison similar to D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1]. We also compared against a simple parallel version of the dual Dijkstra. Each data point is sampled at random meaning a source and destination node is chosen randomly. Each query is run 100 times for all approaches and averaged across all runs. Furthermore, we increase the number of queries in order to measure the throughput of the algorithms. The set of POIs is uniformly sampled from the nodes in the road network with a rate. The rate is multiplied with the total number of nodes in order to get the number of sampled nodes.

4.1.2 Baseline Approach

The dual Dijkstra serves as a baseline for the *in-path* oracle. As a query we used the sampled data points consisting of source and destination pairs.

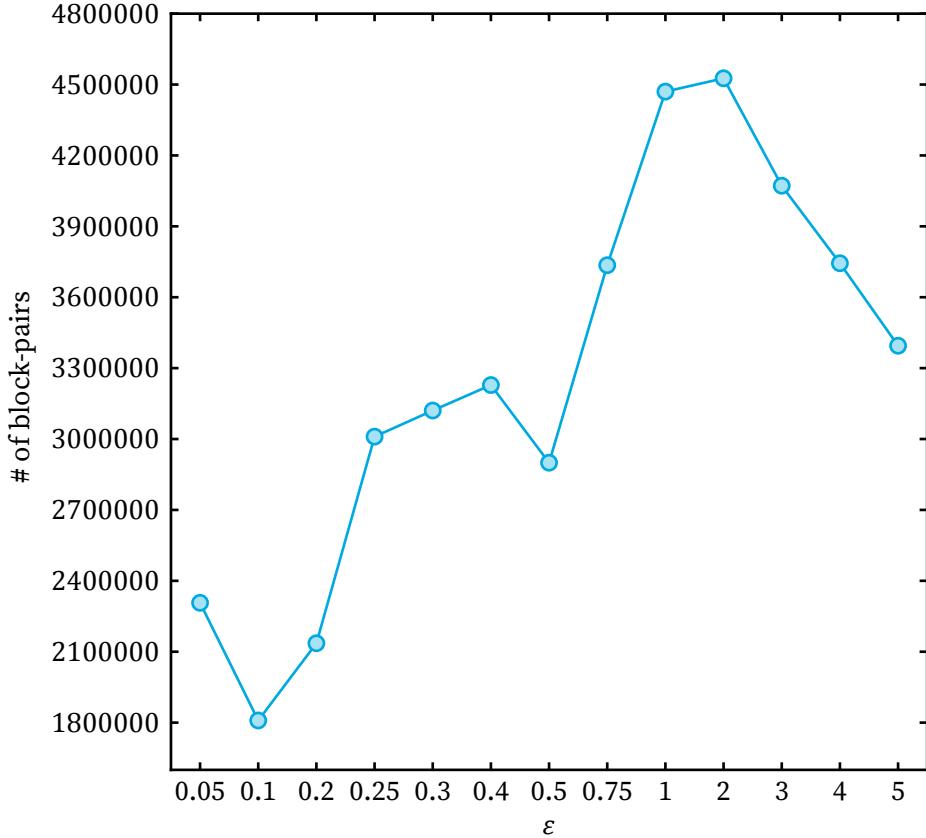
4.2 In-Path Oracle

To measure the performance we examine the size of the oracle with varying the detour limits and road network size as well as the throughput.

4.2.1 Varying Detour Limits

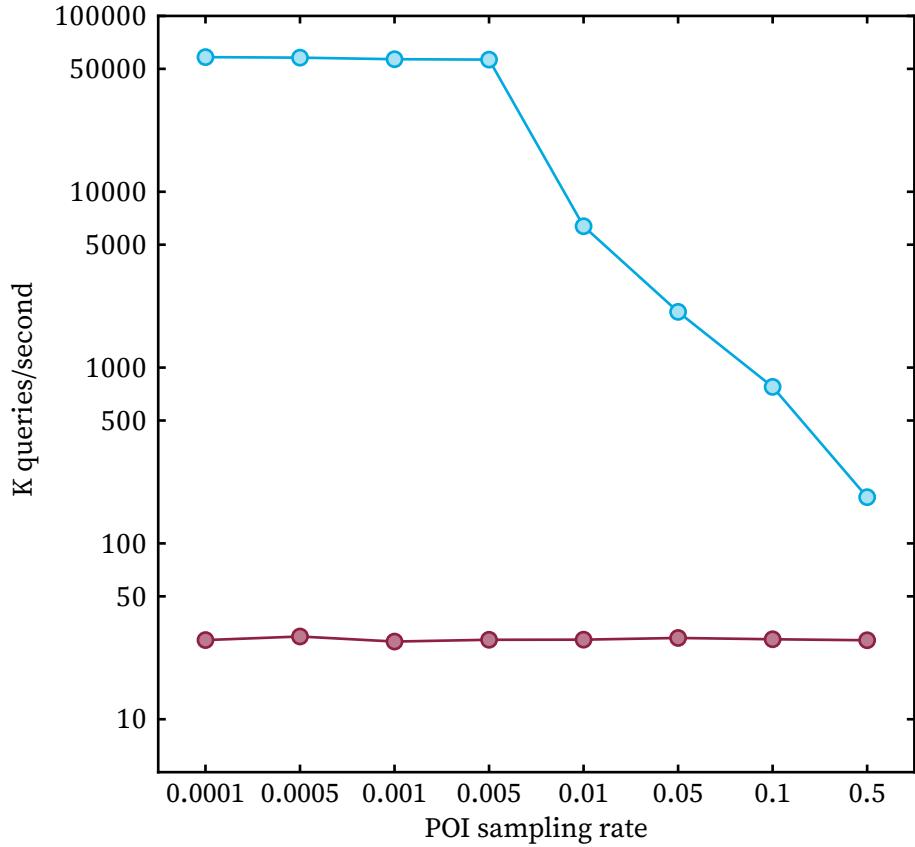
To measure the impact of the detour limit on the oracle size we varied the detour limit from 0.05 to 5. The test were performed on the Konstanz data set consisting of 2282 nodes and 4377 edges. As we can see in Figure 3 the oracle size is roughly shaped like a bell which makes sense when looking at Lemma 3 and Lemma 4. When ε is very small Lemma 4 is more easily satisfied. Similarly when ε is very big Lemma 3 is satisfied for bigger blocks. It is important to note D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1] report much smaller sizes for a graph of this size. For a graph with

5000 nodes they report a oracle size of a little bit more than 100,000 compared to the 3,010,095 (see Figure 3) we found for a graph with 2248 nodes.



4.3 Throughput Experiment

We tested the throughput of *in-path* queries on both the baseline dual Dijkstra and the *in-path* oracle. The experiments were performed on the Konstanz dataset. POIs were randomly sampled with a sampling rate from the dataset which was varied throughout the experiment. We computed the *in-path* oracle for each POI and inserted it into an R*-Tree. Each query was performed on the dual Dijkstra, the parallel dual Dijkstra and the *in-path* oracle.



We observe a constant throughput of about 28,000 *in-path* queries/second for the dual Dijkstra on most POI sampling rates running on only one single thread. The parallel version performed even worse most of the time. However, as expected the *in-path* oracle has a much higher throughput than the dual Dijkstra. Figure 4 clearly shows we get more than 100,000 *in-path* queries per second for all sampling rates. This confirms the findings of D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1].

5 Conclusions and Future Work

We looked at the solution to the *beer-path* problem proposed by D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1] and implemented it in Rust. We could somewhat verify the results with regard to the throughput on small instances. On bigger instances the time to compute the oracle is to long to be practically feasable. The oracle size though we find to be bigger by a factor of more than 10 and also exceeds the upper bound they presented which could be why the compute time is so high. This obviously has an impact on the throughput because of the massive increase in search space (see Figure 4). Because the size of the oracle exceeded the bound presented by D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet [1] further work should be

conducted to provide a concrete proof. Furthermore we find Lemma 3 to be insufficient. Precisely the term $d_N(a_r, b_r) - (r_a^F + r_b^B)$ can be less than 0 because $d_N(a_r, b_r) > (r_a^F + r_b^B)$ is not guaranteed which is why the isolation of ε is not possible. It remains to be seen what the impact of this insufficiency is.

Bibliography

- [1] D. Ghosh, J. Sankaranarayanan, K. Khatter, and H. Samet, “In-Path Oracles for Road Networks,” *ISPRS International Journal of Geo-Information*, vol. 12, no. 7, p. 277, Jul. 2023, doi: 10.3390/ijgi12070277.
- [2] J. Sankaranarayanan, H. Alborzi, and H. Samet, “Efficient query processing on spatial networks,” in *Proceedings of the 13th Annual ACM International Workshop on Geographic Information Systems*, in GIS '05. Bremen, Germany: Association for Computing Machinery, 2005, pp. 200–209. doi: 10.1145/1097064.1097093.
- [3] J. Sankaranarayanan and H. Samet, “Distance Oracles for Spatial Networks,” in *2009 IEEE 25th International Conference on Data Engineering*, IEEE, Mar. 2009, pp. 652–663. doi: 10.1109/icde.2009.53.
- [4] N. Beckmann, H.-P. Kriegel, R. Schneider, and B. Seeger, “The R*-tree: an efficient and robust access method for points and rectangles,” in *Proceedings of the 1990 ACM SIGMOD international conference on Management of data*, in SIGMOD 90. ACM, May 1990. doi: 10.1145/93597.98741.
- [5] A. Guttman, “R-trees: a dynamic index structure for spatial searching,” in *Proceedings of the 1984 ACM SIGMOD international conference on Management of data - SIGMOD '84*, in SIGMOD '84. ACM Press, 1984, p. 47. doi: 10.1145/602259.602266.